

Lab 8

Q1: Find out if following variables are significant or insignificant and need to be dropped.

- i) Seller-insignificant
- ii) offerType-insignificant
- iii) abtest-insignificant
- Iv)vehicleType-significant
- V)gearbox,
- Vi)Model
- Vii)Kilometer
- Viii)Fueltype
- Ix)Brand
- X)notRepairedDamage

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.metrics import recall_score
from sklearn.metrics import precision_score
from sklearn.metrics import f1_score
```

```
In [2]: df = pd.read_csv('cars_sample.csv', encoding = "ISO-8859-1")
```

```
In [3]: df.head()
```

Out[3]:

	dateCrawled	name	seller	offerType	price	abtest	vehicleType	yearOfRegistration	gearbox	powerF
0	30/03/2016 13:51	Zu_verkaufen	private	offer	4450	test	limousine	2003	manual	15
1	7/3/2016 9:54	Volvo_XC90_2.4D_Summum	private	offer	13299	control	suv	2005	manual	16
2	1/4/2016 0:57	Volkswagen_Touran	private	offer	3200	test	bus	2003	manual	10
3	19/03/2016 17:50	Seat_Ibiza_1.4_16V_Reference	private	offer	4500	control	small car	2006	manual	8
4	16/03/2016 14:51	Volvo_XC90_D5_Aut_RDesign_R_Design_AWD_GSHD_S...	private	offer	18750	test	suv	2008	automatic	18

```
In [4]: cols = ['dateCrawled', 'name', 'dateCreated', 'lastSeen']
df.drop(columns=cols, inplace=True)
df.head()
```

Out[4]:

	seller	offerType	price	abtest	vehicleType	yearOfRegistration	gearbox	powerPS	model	kilometer	monthOfRegistration	fuelType	brand	nc
0	private	offer	4450	test	limousine	2003	manual	150	3er	150000	3	diesel	bmw	
1	private	offer	13299	control	suv	2005	manual	163	xc_reihe	150000	6	diesel	volvo	
2	private	offer	3200	test	bus	2003	manual	101	touran	150000	11	diesel	volkswagen	
3	private	offer	4500	control	small car	2006	manual	86	ibiza	60000	12	petrol	seat	
4	private	offer	18750	test	suv	2008	automatic	185	xc_reihe	150000	11	diesel	volvo	

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50001 entries, 0 to 50000
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   seller                 50001 non-null  object
1   offerType              50001 non-null  object
2   price                  50001 non-null  int64
3   abtest                 50001 non-null  object
4   vehicleType            44813 non-null  object
5   yearOfRegistration      50001 non-null  int64
6   gearbox                47177 non-null  object
7   powerPS                50001 non-null  int64
8   model                  47243 non-null  object
9   kilometer              50001 non-null  int64
10  monthOfRegistration     50001 non-null  int64
11  fuelType                45498 non-null  object
12  brand                  50001 non-null  object
13  notRepairedDamage      40285 non-null  object
14  postalCode              50001 non-null  int64
dtypes: int64(6), object(9)
memory usage: 5.7+ MB
```

```
In [6]: for col in df.columns:
print(col)
print(df[col].value_counts())
print()
```

```
seller
private      49999
commercial      2
Name: seller, dtype: int64
```

```
offerType
offer      49998
request      3
Name: offerType, dtype: int64
```

```
price
0      1451
500     742
1500    705
1000     647
2500     594
...
103990    1
370000    1
2151      1
225000    1
175000    1
Name: price, Length: 2393, dtype: int64
```

```
abtest
test      25869
control   24132
Name: abtest, dtype: int64
```

```
vehicleType
limousine      13041
small car      10744
station wagon   8990
bus            3995
cabrio         3056
coupe          2536
suv            2011
others         440
Name: vehicleType, dtype: int64
```

```
yearOfRegistration
2000      3315
2005      3131
1999      3055
2001      2804
2003      2756
...
1928        1
2900        1
8500        1
1940        1
1934        1
Name: yearOfRegistration, Length: 97, dtype: int64
```

```
gearbox
manual      36732
automatic   10445
Name: gearbox, dtype: int64
```

```
powerPS
0      5605
75     3264
60     2167
150    2057
140    1795
...
268        1
416        1
382        1
401        1
386        1
Name: powerPS, Length: 460, dtype: int64
```

```
model
golf      3972
others    3441
3er       2816
polo      1780
corsa     1701
...
serie_2    1
serie_3    1
elefantino 1
charade    1
rangerover 1
Name: model, Length: 248, dtype: int64
```

```
kilometer
150000    32442
125000     5124
100000     2130
90000      1634
80000      1509
70000      1284
60000      1194
50000      1023
5000       1002
40000       856
30000       787
20000       754
10000       262
Name: kilometer, dtype: int64
```

```
monthOfRegistration
0      5043
3      4755
6      4449
4      4153
5      4109
7      3897
10     3666
9      3453
11     3436
12     3403
1      3286
8      3240
2      3111
```

```
2      3111
Name: monthOfRegistration, dtype: int64
```

```
fuelType
petrol      30214
diesel      14347
lpg          778
cng          80
hybrid       39
other        26
electro      14
Name: fuelType, dtype: int64
```

```
brand
volkswagen    10646
bmw            5507
opel           5392
mercedes_benz  4761
audi           4463
ford           3345
renault        2457
peugeot        1513
fiat           1238
seat           1007
mazda          769
skoda          767
nissan          729
smart          718
citroen        698
toyota         602
sonstige_autos  546
volvo          476
hyundai        468
mini           446
mitsubishi     419
honda          347
kia            325
porsche        312
suzuki         301
alfa_romeo     284
chevrolet      244
chrysler       187
dacia          136
subaru         130
```

dacia	136
subaru	130
jeep	106
land_rover	92
jaguar	87
daihatsu	86
trabant	84
saab	74
daewoo	73
lancia	70
rover	63
lada	33

Name: brand, dtype: int64

notRepairedDamage

no	35337
yes	4948

Name: notRepairedDamage, dtype: int64

postalCode

10115	116
65428	70
50354	48
53757	47
44145	47
...	
67715	1
72229	1
74638	1
76764	1
17214	1

Name: postalCode, Length: 7018, dtype: int64


```
In [7]: cols = ['seller', 'offerType', 'notRepairedDamage']
df.drop(columns=cols, inplace=True)
df.head()
```

Out[7]:

	price	abtest	vehicleType	yearOfRegistration	gearbox	powerPS	model	kilometer	monthOfRegistration	fuelType	brand	postalCode
0	4450	test	limousine	2003	manual	150	3er	150000	3	diesel	bmw	20257
1	13299	control	suv	2005	manual	163	xc_reihe	150000	6	diesel	volvo	88045
2	3200	test	bus	2003	manual	101	touran	150000	11	diesel	volkswagen	27449
3	4500	control	small car	2006	manual	86	ibiza	60000	12	petrol	seat	34537
4	18750	test	suv	2008	automatic	185	xc_reihe	150000	11	diesel	volvo	55270

```
In [8]: cols = ['abtest', 'vehicleType', 'gearbox', 'model', 'fuelType', 'brand']
```

```
for col in cols:
    le = preprocessing.LabelEncoder()
    df[col] = le.fit_transform(df[col].astype(str))
df.head()
```

Out[8]:

	price	abtest	vehicleType	yearOfRegistration	gearbox	powerPS	model	kilometer	monthOfRegistration	fuelType	brand	postalCode
0	4450	1	3	2003	1	150	11	150000	3	1	2	20257
1	13299	0	8	2005	1	163	243	150000	6	1	39	88045
2	3200	1	0	2003	1	101	221	150000	11	1	38	27449
3	4500	0	6	2006	1	86	120	60000	12	7	30	34537
4	18750	1	8	2008	0	185	243	150000	11	1	39	55270

```
In [9]: df.shape
```

Out[9]: (50001, 12)

```
In [10]: df.corr()
```

Out[10]:

	price	abtest	vehicleType	yearOfRegistration	gearbox	powerPS	model	kilometer	monthOfRegistration	fuelType	brand	postalCode
price	1.000000	0.002790	-0.011208	0.017604	-0.018165	0.020429	-0.002403	-0.045458	0.000582	-0.013127	-0.007697	0.005916
abtest	0.002790	1.000000	0.005034	0.003324	-0.003996	0.001375	-0.001415	-0.003027	0.000621	0.004686	0.006246	0.003096
vehicleType	-0.011208	0.005034	1.000000	0.000573	-0.000225	-0.035590	-0.037315	0.020446	0.006124	-0.035184	0.012066	-0.013254
yearOfRegistration	0.017604	0.003324	0.000573	1.000000	0.029205	-0.004394	0.008299	-0.064188	-0.023152	-0.012598	0.004461	-0.001615
gearbox	-0.018165	-0.003996	-0.000225	0.029205	1.000000	-0.142459	0.046735	0.005481	-0.123792	0.126904	0.120576	0.003200
powerPS	0.020429	0.001375	-0.035590	-0.004394	-0.142459	1.000000	-0.035191	-0.016447	0.034345	-0.044093	-0.083801	0.017415
model	-0.002403	-0.001415	-0.037315	0.008299	0.046735	-0.035191	1.000000	-0.043010	-0.028372	-0.034566	0.435585	-0.051870
kilometer	-0.045458	-0.003027	0.020446	-0.064188	0.005481	-0.016447	-0.043010	1.000000	0.001985	-0.104424	-0.104424	-0.024076
monthOfRegistration	0.000582	0.000621	0.006124	-0.023152	-0.123792	0.034345	-0.028372	0.001985	1.000000	-0.062377	0.038798	0.019050
fuelType	-0.013127	0.004686	-0.035184	-0.012598	0.126904	-0.044093	-0.034566	-0.104424	-0.062377	1.000000	0.038798	-0.014061
brand	-0.007697	0.006246	0.012066	0.004461	0.120576	-0.083801	0.435585	-0.031284	-0.018635	0.038798	1.000000	-0.014061
postalCode	0.005916	0.003096	-0.013254	-0.001615	0.003200	0.017415	-0.051870	-0.024076	0.019050	-0.014061	-0.014061	1.000000

Q2: Drop insignificant variables from dataframe 'cars'.

```
In [11]: ▶ cols = ['abtest', 'model', 'monthOfRegistration', 'brand', 'postalCode']  
df.drop(columns=cols, inplace=True)  
df.head()
```


Out[11]:

	price	vehicleType	yearOfRegistration	gearbox	powerPS	kilometer	fuelType
0	4450	3	2003	1	150	150000	1
1	13299	8	2005	1	163	150000	1
2	3200	0	2003	1	101	150000	1
3	4500	6	2006	1	86	60000	7
4	18750	8	2008	0	185	150000	1

Q3: Find correlation between all numerical variables and find which variable has the highest correlation with price.

```
In [12]: cor = df.corr()
```

```
In [13]: sns.heatmap(cor, cmap="YlGnBu", annot=True)  
plt.show()
```



```
In [14]: print(cor['price'])  
print('Highest: kilometer (abs vale of 0.045458)')
```

```
price          1.000000  
vehicleType   -0.011208  
yearOfRegistration  0.017604  
gearbox       -0.018165  
powerPS        0.020429  
kilometer     -0.045458  
fuelType      -0.013127  
Name: price, dtype: float64  
Highest: kilometer (abs vale of 0.045458)
```

Q4: Calculate the training data and testing data score using a linear regression model.

```
In [15]: x_train, x_test, y_train, y_test = train_test_split(df.drop(columns = ['price']), df['price'], test_size = 0.2)
          x_train.shape, y_train.shape, x_test.shape, y_test.shape
```

```
Out[15]: ((40000, 6), (40000,), (10001, 6), (10001,))
```

```
In [16]: algo = "Linear Regression\n"
          model = LinearRegression()
          model.fit(x_train, y_train)
          print(algo)

          print('Training error')
          y_pred = model.predict(x_train)
          e = (y_pred - y_train)
          e = e.dot(e)
          e /= y_test.shape[0]
          e = e**0.5
          print(e)

          print('Testing error')
          y_pred = model.predict(x_test)
          e = (y_pred - y_test)
          e = e.dot(e)
          e /= y_test.shape[0]
          e = e**0.5
          print(e)
```

Linear Regression

Training error
160444.06822035677
Testing error
104703.69483733781