

# Lab 9

**Q1:** From dataset 'social ad':  
Calculate performance metric

1. Accuracy
2. Sensitivity/Recall
3. Specificity
4. F-score
5. Precision

Use Algorithmk-NN

1. Logistic regression
2. SVM
3. Decision Tree
4. Naive Bayes

Plot a bar graph and compare the accuracy obtained in each case.

CODE:

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn import preprocessing
import seaborn as sns

from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.metrics import recall_score
from sklearn.metrics import precision_score
from sklearn.metrics import f1_score
```

```
In [2]: import warnings
warnings.filterwarnings('ignore')
```

```
In [3]: df = pd.read_csv('Social_Network_Ads.csv')
```

```
In [4]: df.head()
```

```
Out[4]:
```

	Age	EstimatedSalary	Purchased
0	19	19000	0
1	35	20000	0
2	26	43000	0
3	27	57000	0
4	19	76000	0

In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   Age              400 non-null   int64
1   EstimatedSalary  400 non-null   int64
2   Purchased        400 non-null   int64
dtypes: int64(3)
memory usage: 9.5 KB
```

In [6]: `corr = df.corr()`  
`corr.style.background_gradient(cmap='coolwarm')`

Out[6]:

	Age	EstimatedSalary	Purchased
Age	1.000000	0.155238	0.622454
EstimatedSalary	0.155238	1.000000	0.362083
Purchased	0.622454	0.362083	1.000000

In [7]: `x_train, x_test, y_train, y_test = train_test_split(df.drop(columns = ['Purchased']), df['Purchased'], test_size = 0.2)`  
`x_train.shape, y_train.shape, x_test.shape, y_test.shape`

Out[7]: `((320, 2), (320,), (80, 2), (80,))`

In [8]: `algos = []`  
`accuracy = []`  
`recall = []`  
`precision = []`  
`f1Score = []`  
`specificity = []`

```
In [9]: ► algo = "Logistic Regression"
model = LogisticRegression()
model.fit(x_train, y_train)
y_pred = model.predict(x_test)
print(algo)
print(confusion_matrix(y_test, y_pred), '\n\n')
acc = accuracy_score(y_test, y_pred) * 100
print('Accuracy:', acc)
rec = recall_score(y_test, y_pred) * 100
print('Recall:', rec)
pre = precision_score(y_test, y_pred) * 100
print('Precision:', pre)
f1s = f1_score(y_test, y_pred) * 100
print('F score:', f1s)
tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
spc = tn / (tn+fp)
print('Specificity:', spc)

algos.append(algo)
accuracy.append(acc)
recall.append(rec)
precision.append(pre)
f1Score.append(f1s)
specificity.append(spc)
```

```
Logistic Regression
[[45  0]
 [35  0]]
```

```
Accuracy: 56.25
Recall: 0.0
Precision: 0.0
F score: 0.0
Specificity: 1.0
```

```
In [10]: ► algo = "K Nearest Neighbour"
model = KNeighborsClassifier()
model.fit(x_train, y_train)
y_pred = model.predict(x_test)
print(algo)
print(confusion_matrix(y_test, y_pred), '\n\n')
acc = accuracy_score(y_test, y_pred) * 100
print('Accuracy:', acc)
rec = recall_score(y_test, y_pred) * 100
print('Recall:', rec)
pre = precision_score(y_test, y_pred) * 100
print('Precision:', pre)
f1s = f1_score(y_test, y_pred) * 100
print('F score:', f1s)
tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
spc = tn / (tn+fp)
print('Specificity:', spc)

algos.append(algo)
accuracy.append(acc)
recall.append(rec)
precision.append(pre)
f1Score.append(f1s)
specificity.append(spc)

K Nearest Neighbour
[[37  8]
 [13 22]]

Accuracy: 73.75
Recall: 62.857142857142854
Precision: 73.33333333333333
F score: 67.69230769230768
Specificity: 0.8222222222222222
```

```
In [11]: ▶ algo = "Decision Tree"
model = DecisionTreeClassifier()
model.fit(x_train, y_train)
y_pred = model.predict(x_test)
print(algo)
print(confusion_matrix(y_test, y_pred), '\n\n')
acc = accuracy_score(y_test, y_pred) * 100
print('Accuracy:', acc)
rec = recall_score(y_test, y_pred) * 100
print('Recall:', rec)
pre = precision_score(y_test, y_pred) * 100
print('Precision:', pre)
f1s = f1_score(y_test, y_pred) * 100
print('F score:', f1s)
tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
spc = tn / (tn+fp)
print('Specificity:', spc)

algos.append(algo)
accuracy.append(acc)
recall.append(rec)
precision.append(pre)
f1Score.append(f1s)
specificity.append(spc)
```

Decision Tree

```
[[41  4]
 [ 8 27]]
```

Accuracy: 85.0

Recall: 77.14285714285715

Precision: 87.09677419354838

F score: 81.81818181818183

Specificity: 0.9111111111111111



```
In [12]: ► algo = "Naive Bayes"
model = GaussianNB()
model.fit(x_train, y_train)
y_pred = model.predict(x_test)
print(algo)
print(confusion_matrix(y_test, y_pred), '\n\n')
acc = accuracy_score(y_test, y_pred) * 100
print('Accuracy:', acc)
rec = recall_score(y_test, y_pred) * 100
print('Recall:', rec)
pre = precision_score(y_test, y_pred) * 100
print('Precision:', pre)
f1s = f1_score(y_test, y_pred) * 100
print('F score:', f1s)
tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
spc = tn / (tn+fp)
print('Specificity:', spc)

algos.append(algo)
accuracy.append(acc)
recall.append(rec)
precision.append(pre)
f1Score.append(f1s)
specificity.append(spc)
```

Naive Bayes

```
[[42  3]
 [13 22]]
```

Accuracy: 80.0

Recall: 62.857142857142854

Precision: 88.0

F score: 73.33333333333334

Specificity: 0.9333333333333333

```
In [13]: ► algo = "SVM"
model = SVC(kernel='rbf')
model.fit(x_train, y_train)
y_pred = model.predict(x_test)
print(algo)
print(confusion_matrix(y_test, y_pred), '\n\n')
acc = accuracy_score(y_test, y_pred) * 100
print('Accuracy:', acc)
rec = recall_score(y_test, y_pred) * 100
print('Recall:', rec)
pre = precision_score(y_test, y_pred) * 100
print('Precision:', pre)
f1s = f1_score(y_test, y_pred) * 100
print('F score:', f1s)
tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
spc = tn / (tn+fp)
print('Specificity:', spc)

algos.append(algo)
accuracy.append(acc)
recall.append(rec)
precision.append(pre)
f1Score.append(f1s)
specificity.append(spc)
```

```
SVM
[[43  2]
 [22 13]]
```

```
Accuracy: 70.0
Recall: 37.142857142857146
Precision: 86.66666666666667
F score: 52.0
Specificity: 0.9555555555555556
```

```
In [14]: ▶ algo = "Adaboost"
model = AdaBoostClassifier()
model.fit(x_train, y_train)
y_pred = model.predict(x_test)
print(algo)
print(confusion_matrix(y_test, y_pred), '\n\n')
acc = accuracy_score(y_test, y_pred) * 100
print('Accuracy:', acc)
rec = recall_score(y_test, y_pred) * 100
print('Recall:', rec)
pre = precision_score(y_test, y_pred) * 100
print('Precision:', pre)
f1s = f1_score(y_test, y_pred) * 100
print('F score:', f1s)
tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
spc = tn / (tn+fp)
print('Specificity:', spc)

algos.append(algo)
accuracy.append(acc)
recall.append(rec)
precision.append(pre)
f1Score.append(f1s)
specificity.append(spc)
```

```
Adaboost
[[40  5]
 [ 9 26]]
```

```
Accuracy: 82.5
Recall: 74.28571428571429
Precision: 83.87096774193549
F score: 78.7878787878788
Specificity: 0.8888888888888888
```



```
In [15]: ▶ algo = "GradientBoost"
model = GradientBoostingClassifier()
model.fit(x_train, y_train)
y_pred = model.predict(x_test)
print(algo)
print(confusion_matrix(y_test, y_pred), '\n\n')
acc = accuracy_score(y_test, y_pred) * 100
print('Accuracy:', acc)
rec = recall_score(y_test, y_pred) * 100
print('Recall:', rec)
pre = precision_score(y_test, y_pred) * 100
print('Precision:', pre)
f1s = f1_score(y_test, y_pred) * 100
print('F score:', f1s)
tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
spc = tn / (tn+fp)
print('Specificity:', spc)

algorithms.append(algo)
accuracy.append(acc)
recall.append(rec)
precision.append(pre)
f1Score.append(f1s)
specificity.append(spc)
```

GradientBoost

```
[[40  5]
 [ 7 28]]
```

Accuracy: 85.0

Recall: 80.0

Precision: 84.84848484848484

F score: 82.3529411764706

Specificity: 0.8888888888888888

```
In [16]: ► algo = "Random forest"
model = RandomForestClassifier()
model.fit(x_train, y_train)
y_pred = model.predict(x_test)
print(algo)
print(confusion_matrix(y_test, y_pred), '\n\n')
acc = accuracy_score(y_test, y_pred) * 100
print('Accuracy:', acc)
rec = recall_score(y_test, y_pred) * 100
print('Recall:', rec)
pre = precision_score(y_test, y_pred) * 100
print('Precision:', pre)
f1s = f1_score(y_test, y_pred) * 100
print('F score:', f1s)
tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
spc = tn / (tn+fp)
print('Specificity:', spc)

algos.append(algo)
accuracy.append(acc)
recall.append(rec)
precision.append(pre)
f1Score.append(f1s)
specificity.append(spc)
```

```
Random forest
[[40  5]
 [ 6 29]]
```

```
Accuracy: 86.25
Recall: 82.85714285714286
Precision: 85.29411764705883
F score: 84.05797101449276
Specificity: 0.8888888888888888
```

Pooja Agarwal  
1905330

```
In [17]: mx = 0
for i in range(len(algos)):
    print(algos[i], ': ', accuracy[i], ', ', recall[i], ', ', precision[i], ', ', f1score[i])
    if accuracy[i] > accuracy[mx]:
        mx = i

Logistic Regression :      56.25 ,      0.0 ,      0.0 ,      0.0
K Nearest Neighbour :      73.75 ,      62.857142857142854 ,      73.33333333333333 ,      67.69230769230768
Decision Tree :          85.0 ,      77.14285714285715 ,      87.09677419354838 ,      81.81818181818183
Naive Bayes :           80.0 ,      62.857142857142854 ,      88.0 ,      73.33333333333334
SVM :                   70.0 ,      37.142857142857146 ,      86.66666666666667 ,      52.0
Adaboost :              82.5 ,      74.28571428571429 ,      83.87096774193549 ,      78.7878787878788
GradientBoost :         85.0 ,      80.0 ,      84.84848484848484 ,      82.3529411764706
Random forest :         86.25 ,      82.85714285714286 ,      85.29411764705883 ,      84.05797101449276
```

```
In [18]: print('Maximum Accuracy : ', accuracy[i], 'of', algos[i], 'algorithm.')

Maximum Accuracy : 86.25 of Random forest algorithm.
```

```
In [19]: plt.bar(algos, accuracy)
plt.xticks(rotation = 45)
plt.show()
```

